

MicroRNA prediction from genome-wide data with deep learning: a novel approach based on convolutional residual networks

C. Yones, L.A. Bugnon, J. Raad, D.H. Milone and G. Stegmayer
Research Institute for Signals, Systems and Computational Intelligence



1 Introduction

- Background
- Motivation

2 Proposed method

- Convolutional Neural Networks
- Proposed architecture

3 Results

- Experimental setup
- Precision-recall curves

4 Conclusions

1 Introduction

- Background
- Motivation

2 Proposed method

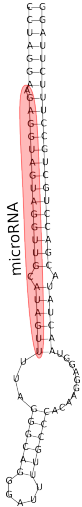
- Convolutional Neural Networks
- Proposed architecture

3 Results

- Experimental setup
- Precision-recall curves

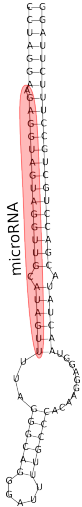
4 Conclusions

Background



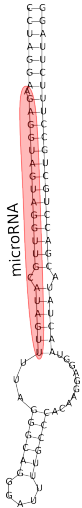
- MicroRNAs (miRNAs) play an essential role in post-transcriptional gene regulation.
- Precursors of miRNA (pre-miRNAs) are characterized by hairpins structure.

Background



- MicroRNAs (miRNAs) play an essential role in post-transcriptional gene regulation.
- Precursors of miRNA (pre-miRNAs) are characterized by hairpins structure.
- A large amount of similar sequences can be folded into this kind of structure.

Background



- MicroRNAs (miRNAs) play an essential role in post-transcriptional gene regulation.
- Precursors of miRNA (pre-miRNAs) are characterized by hairpins structure.
- A large amount of similar sequences can be folded into this kind of structure.
- Machine learning algorithms have been proposed to predict which sequences are likely to contain a miRNA.

Background



- MicroRNAs (miRNAs) play an essential role in post-transcriptional gene regulation.
- Precursors of miRNA (pre-miRNAs) are characterized by hairpins structure.
- A large amount of similar sequences can be folded into this kind of structure.
- Machine learning algorithms have been proposed to predict which sequences are likely to contain a miRNA.

But, there are some problems with ML methods

- Datasets used are not representative of the wide variety of negative examples.
→ Use all hairpins of the genome for validation.
- The performance measures used underestimate the effect of imbalance.
→ Take into account the number of false positives.



But, there are some problems with ML methods



- Datasets used are not representative of the wide variety of negative examples.
→ Use all hairpins of the genome for validation.
- The performance measures used underestimate the effect of imbalance.
→ Take into account the number of false positives.
- The validation methodology does not imitate a real prediction task.
→ Test on sequences from unseen species.

But, there are some problems with ML methods



- Datasets used are not representative of the wide variety of negative examples.
→ Use all hairpins of the genome for validation.
- The performance measures used underestimate the effect of imbalance.
→ Take into account the number of false positives.
- The validation methodology does not imitate a real prediction task.
→ Test on sequences from unseen species.

Taking into account these points, simple sequence alignment works better than machine learning methods.

But, there are some problems with ML methods



- Datasets used are not representative of the wide variety of negative examples.
→ Use all hairpins of the genome for validation.
- The performance measures used underestimate the effect of imbalance.
→ Take into account the number of false positives.
- The validation methodology does not imitate a real prediction task.
→ Test on sequences from unseen species.

Taking into account these points, simple sequence alignment works better than machine learning methods.

Can it be done better?

> cel-mir-62

GUGAGUUAGAUCUCAUAUCCUCCGCAAAAUGGAAAUGAUAUGUAAUCUAGCUUACAG

((((((((((((..((((((..((((((.....))))))..))))))..))))))..))..

Can it be done better?

> cel-mir-62

GUGAGUUAGAUCUCAUAUCCUUCGCAAAAUGGAAAUGAUAUGUAAUCUAGCUUACAG

((((((((((((..((((((..((((((.....))))))..))))))..))))))..))..

> Unknown sequence 1

UUGAGCUAGAUUACAUAUCCGUAAGUGGGGCUUGAAUGAUUAUGUAAUCUAGCUUACAG

.((((((((((((((((((((((((..((((((.....))))))..))))))..))))))..))..

Can it be done better?

> cel-mir-62

GUGAGUUAGAUCUCAUAUCCUCCGCAAAAUGGAAAUGAUAUGUAAUCUAGCUUACAG

((((((((((((..((((((..((((((.....))))))..))))))..))))))..))..

> Unknown sequence 1

UUGAGCUAGAUUACAUAUCCGUAAGUGGGGCUUGAAUGAUAUGUAAUCUAGCUUACAG

.((((((((((((((((((((..((((((...))))))..))))))..))))))..))..

> Unknown sequence 2

GUGAGUUAGAUCUCAUAUCCUCCGCAAAAUGGAAAUCCGAUCUAAUCUACCUUACAG

((((((((((((((..((..((((.....))))..))..))))))..))))..))..

Can it be done better?

> cel-mir-62

GUGAGUUAGAUCUCAUAUCCUCCGCAAAAUGGAAAUGAUAUGUAAUCUAGCUUACAG

((((((((((((..((((((..((((((.....))))))..))))))..))))))..))))))..

> Unknown sequence 1

UUGAGCUAGAUUACAUAUCCGUAAGUGGGGCUUGAAUGAUAUGUAAUCUAGCUUACAG

.((((((((((((((((((((..((((((.....))))))..))))))..))))))..))))))..

Matching score: 189 (clustalw2)

> Unknown sequence 2

GUGAGUUAGAUCUCAUAUCCUCCGCAAAAUGGAAAUCCGAUCUAAUCUACCUUACAG

((((((((((((((..((..((((.....))))..))..))))))..))))..))))..

Matching score: 249 (clustalw2)

1 Introduction

- Background
- Motivation

2 Proposed method

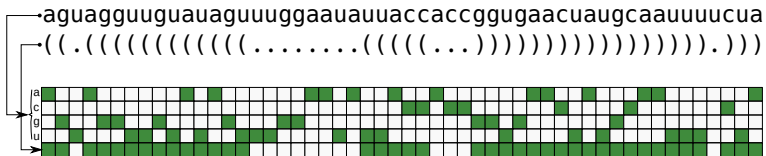
- Convolutional Neural Networks
- Proposed architecture

3 Results

- Experimental setup
- Precision-recall curves

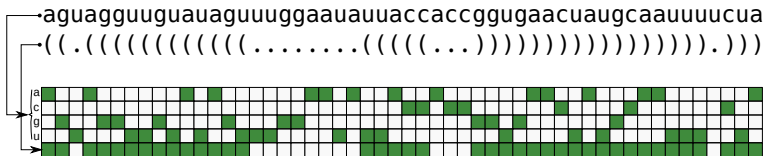
4 Conclusions

A gentle introduction to Convolutional Neural Networks



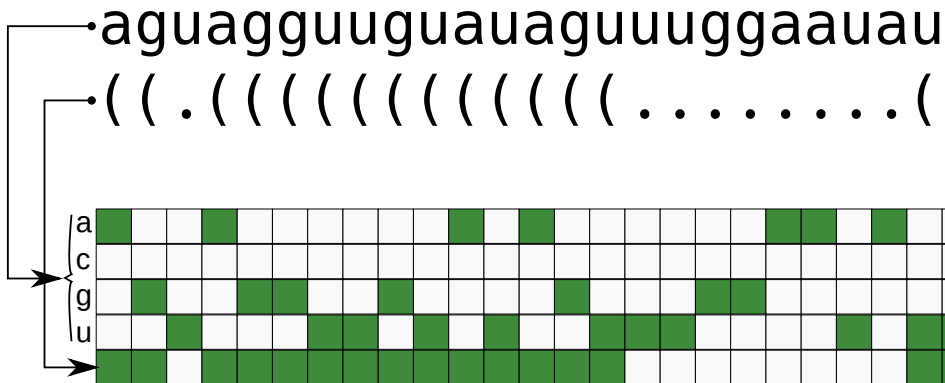
Each column represents a nucleotide

A gentle introduction to Convolutional Neural Networks

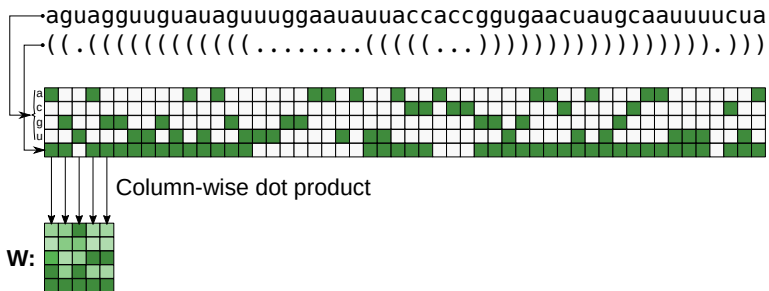


Each column represents a nucleotide

A gentle introduction to Convolutional Neural Networks

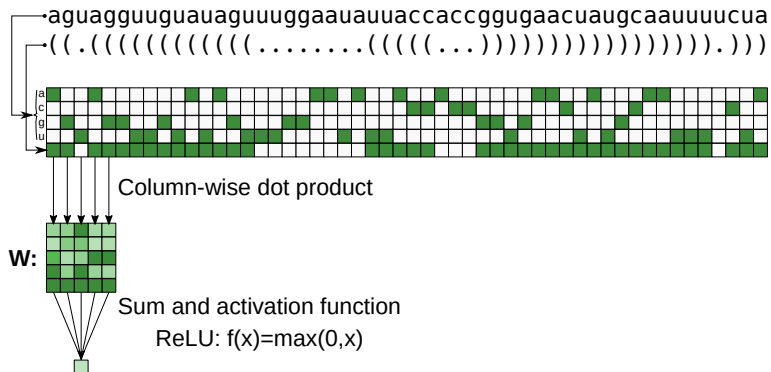


A gentle introduction to Convolutional Neural Networks



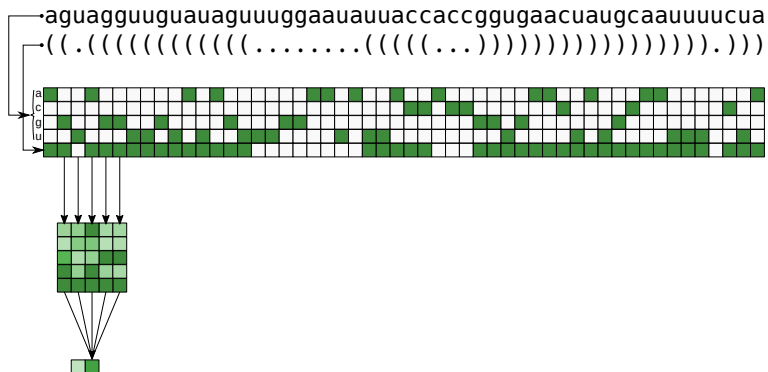
Convolution of the input tensor with the filter **W**

A gentle introduction to Convolutional Neural Networks



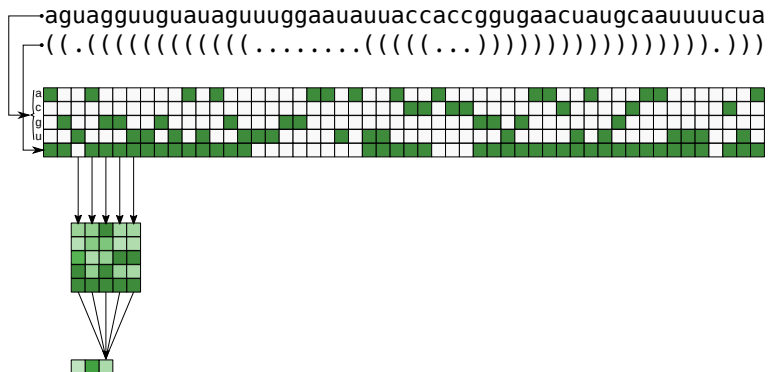
Convolution of the input tensor with the filter **W**

A gentle introduction to Convolutional Neural Networks



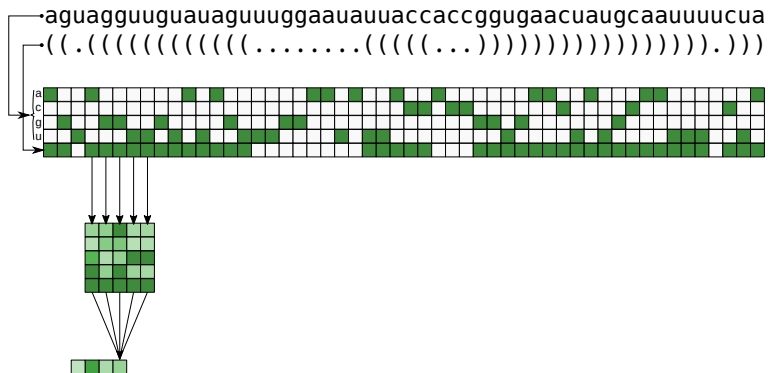
Convolution of the input tensor with the filter W

A gentle introduction to Convolutional Neural Networks



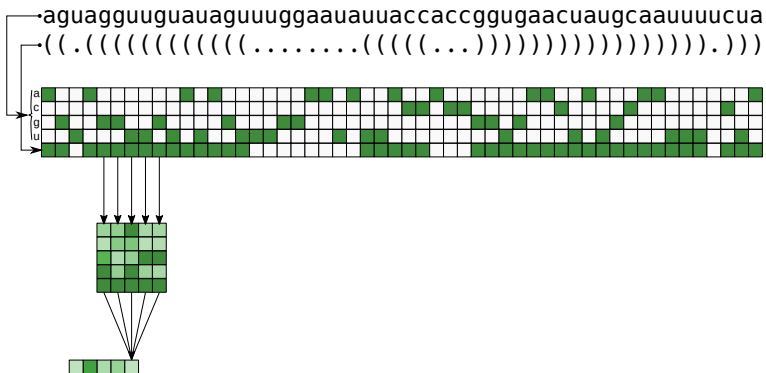
Convolution of the input tensor with the filter **W**

A gentle introduction to Convolutional Neural Networks



Convolution of the input tensor with the filter \mathbf{W}

A gentle introduction to Convolutional Neural Networks



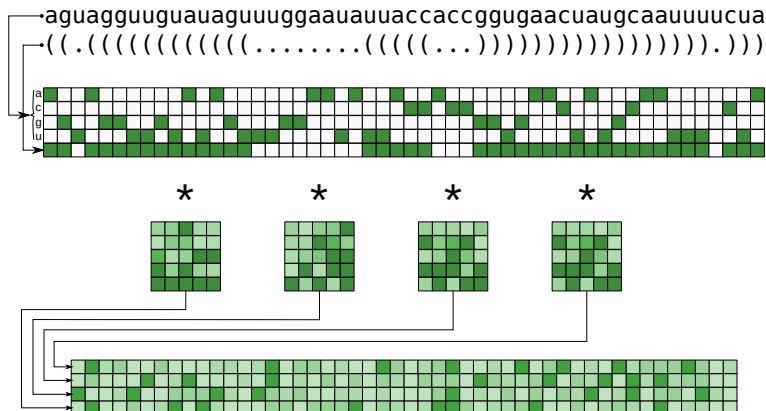
Convolution of the input tensor with the filter **W**

A gentle introduction to Convolutional Neural Networks



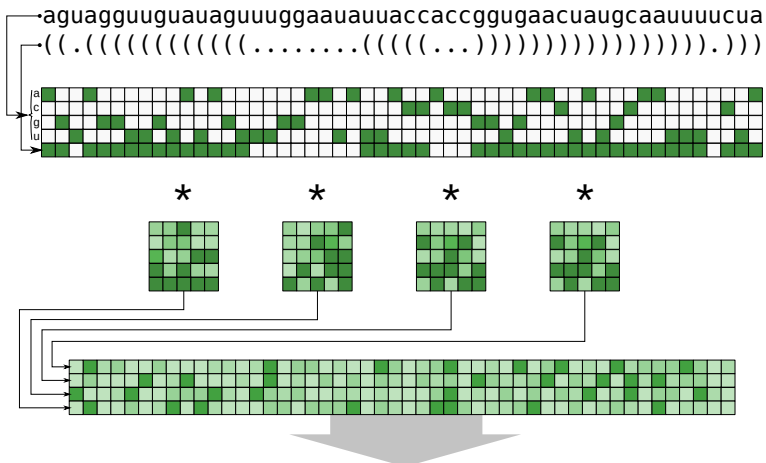
Convolution of the input tensor with the filter **W**

A gentle introduction to Convolutional Neural Networks



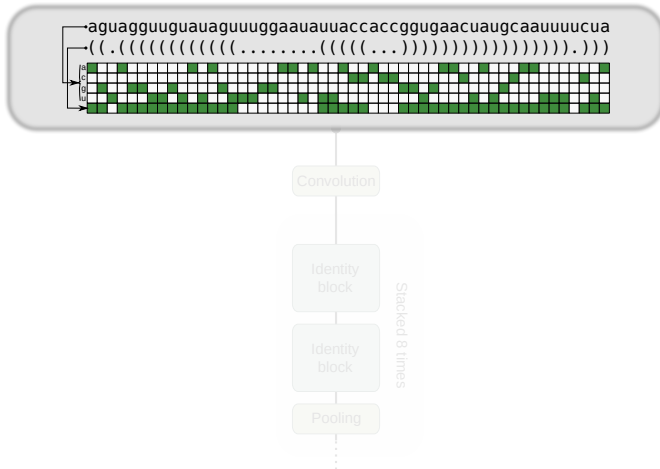
Each filter generates a row in the output tensor

A gentle introduction to Convolutional Neural Networks

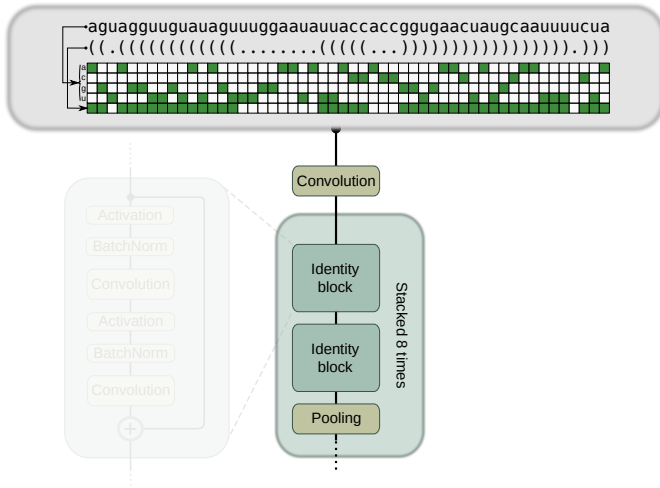


More convolutions and non-linear functions

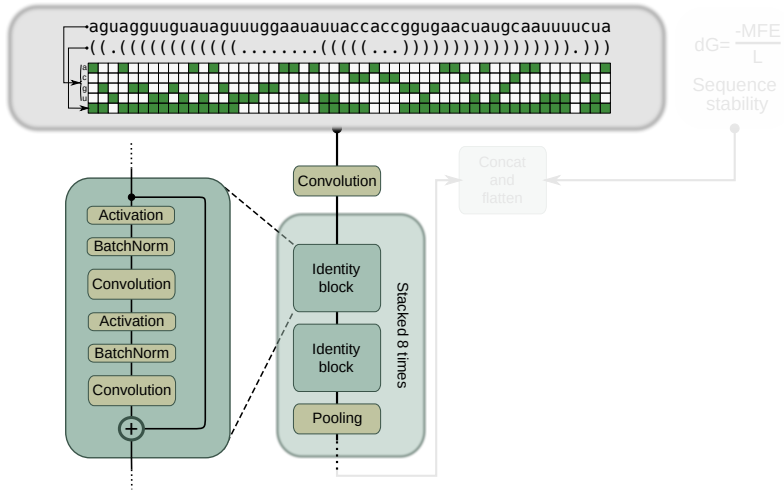
Proposed architecture: mirDNN



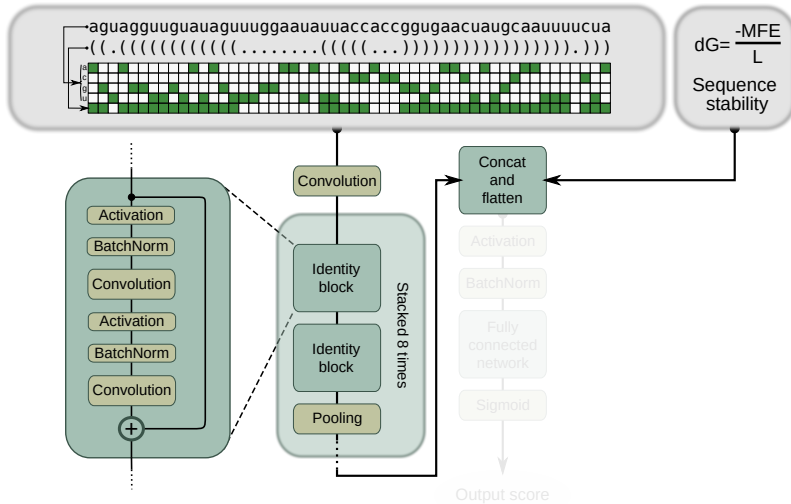
Proposed architecture: mirDNN



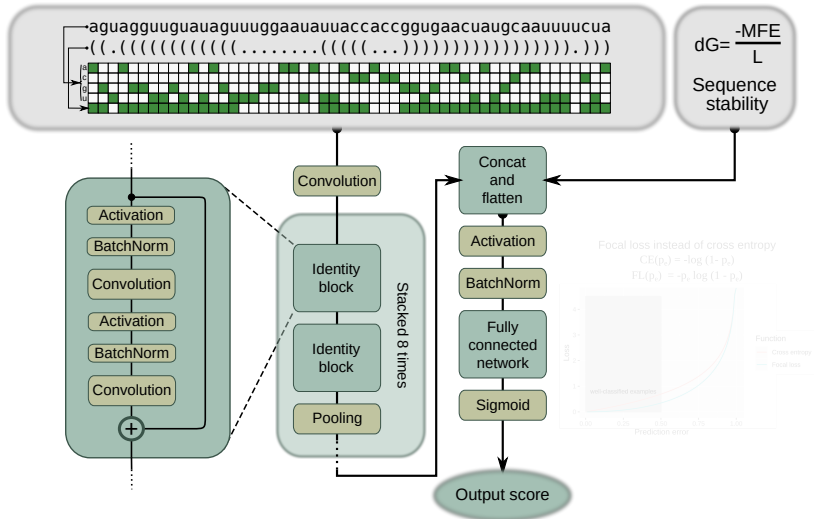
Proposed architecture: mirDNN



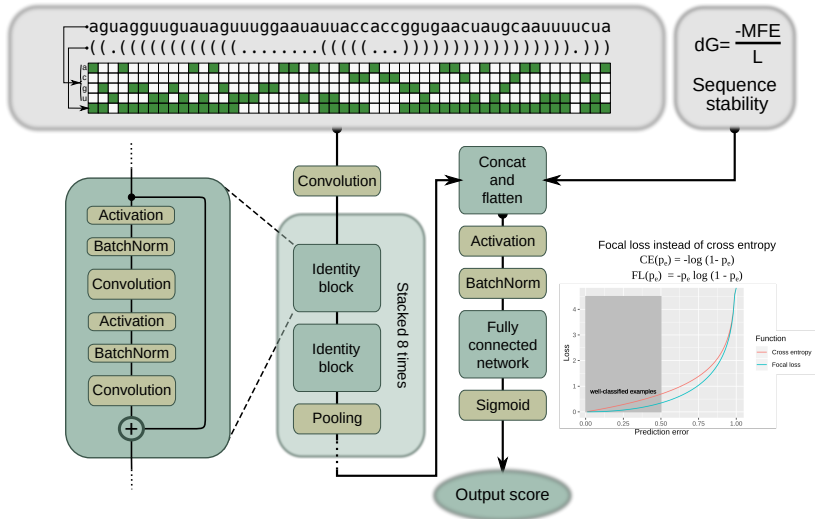
Proposed architecture: mirDNN



Proposed architecture: mirDNN



Proposed architecture: mirDNN

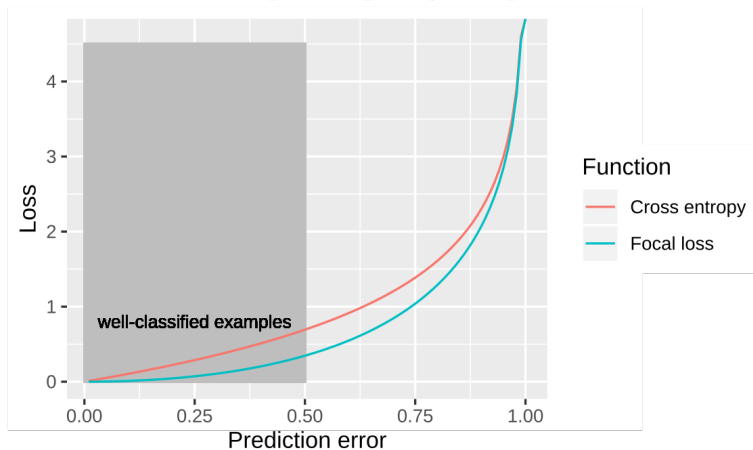


Proposed architecture: mirDNN

Focal loss instead of cross entropy

$$CE(p_e) = -\log(1 - p_e)$$

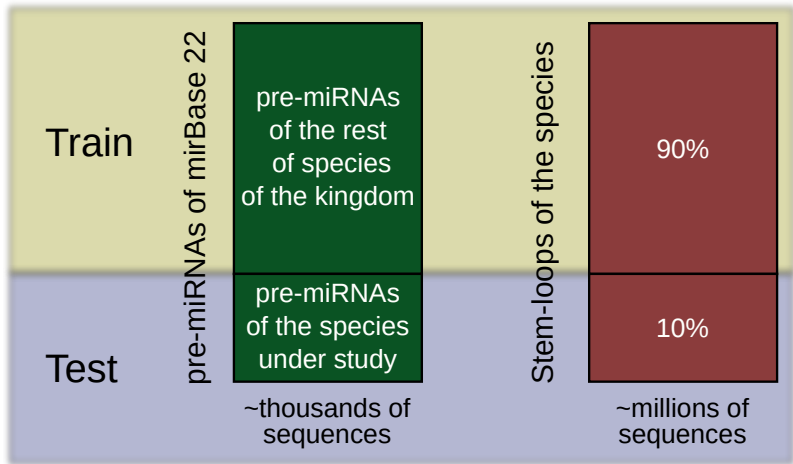
$$FL(p_e) = -p_e \log(1 - p_e)$$



- 1 Introduction
 - Background
 - Motivation
- 2 Proposed method
 - Convolutional Neural Networks
 - Proposed architecture
- 3 Results
 - Experimental setup
 - Precision-recall curves
- 4 Conclusions

Experimental setup

Validation on a *leave-species-out* scheme



Experimental setup

- Test on three well-known species: *Arabidopsis thaliana*, *Caenorhabditis elegans* and *Anopheles gambiae*.
- Two machine learning methods and one sequence alignment method were used for comparison.

Experimental setup

- Test on three well-known species: *Arabidopsis thaliana*, *Caenorhabditis elegans* and *Anopheles gambiae*.
- Two machine learning methods and one sequence alignment method were used for comparison.
- The precision and recall were used as performance measures:

$$Pr = \frac{TP}{TP+FP} \quad Rc = \frac{TP}{TP+FN}$$

Experimental setup

- Test on three well-known species: *Arabidopsis thaliana*, *Caenorhabditis elegans* and *Anopheles gambiae*.
- Two machine learning methods and one sequence alignment method were used for comparison.
- The precision and recall were used as performance measures:

$$Pr = \frac{TP}{TP+FP} \quad Rc = \frac{TP}{TP+FN}$$

- Varying the threshold that defines what is classified as positive or negative, precision-recall curves were generated.

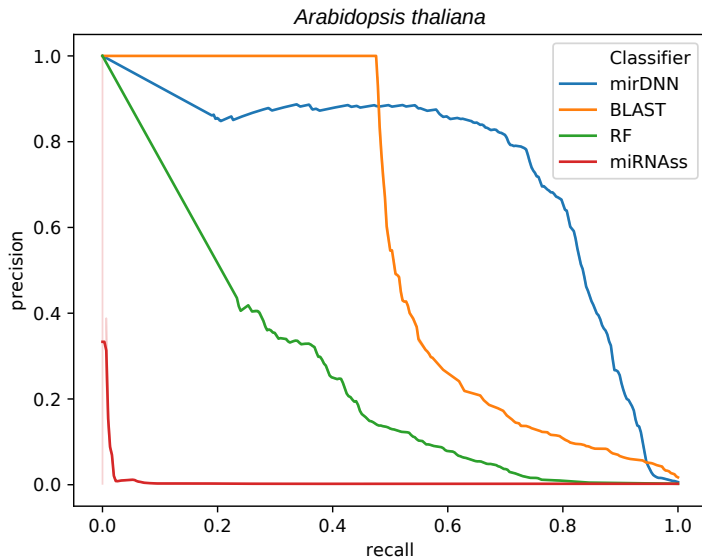
Experimental setup

- Test on three well-known species: *Arabidopsis thaliana*, *Caenorhabditis elegans* and *Anopheles gambiae*.
- Two machine learning methods and one sequence alignment method were used for comparison.
- The precision and recall were used as performance measures:

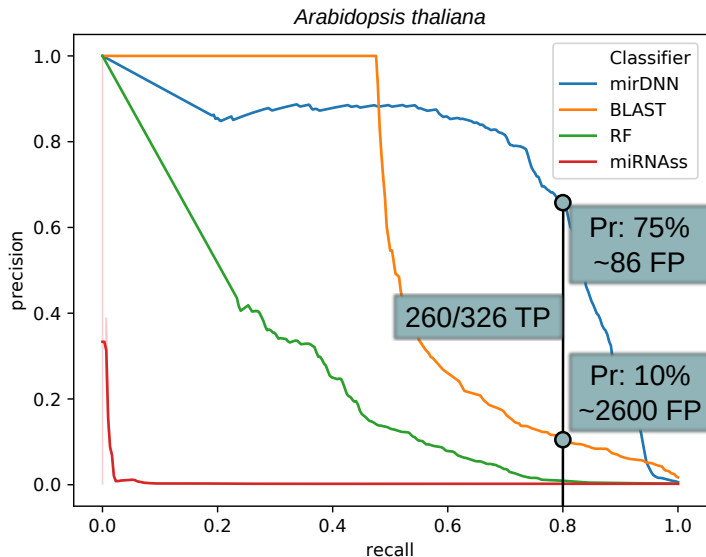
$$Pr = \frac{TP}{TP+FP} \quad Rc = \frac{TP}{TP+FN}$$

- Varying the threshold that defines what is classified as positive or negative, precision-recall curves were generated.

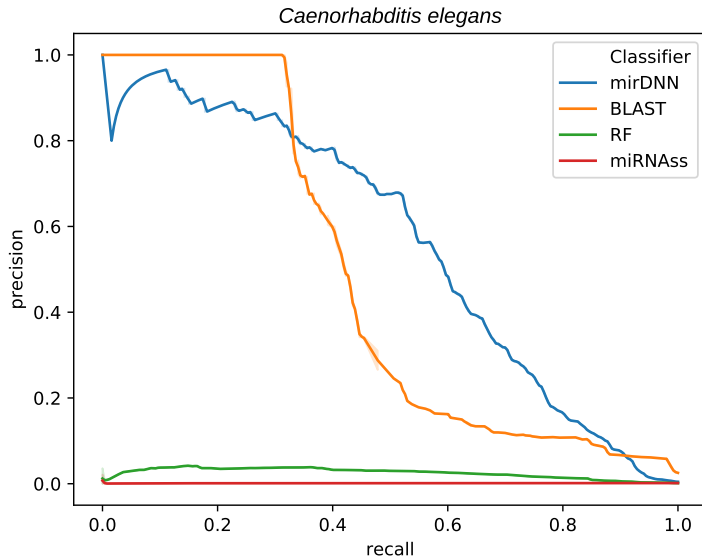
Precision-recall curves



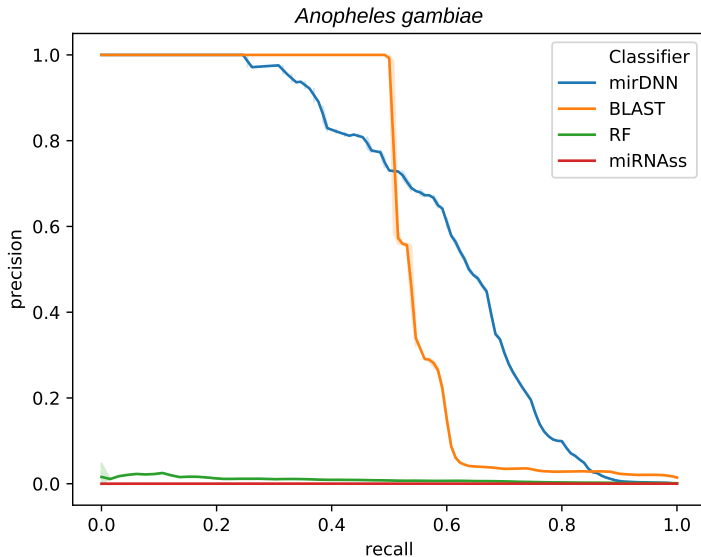
Precision-recall curves



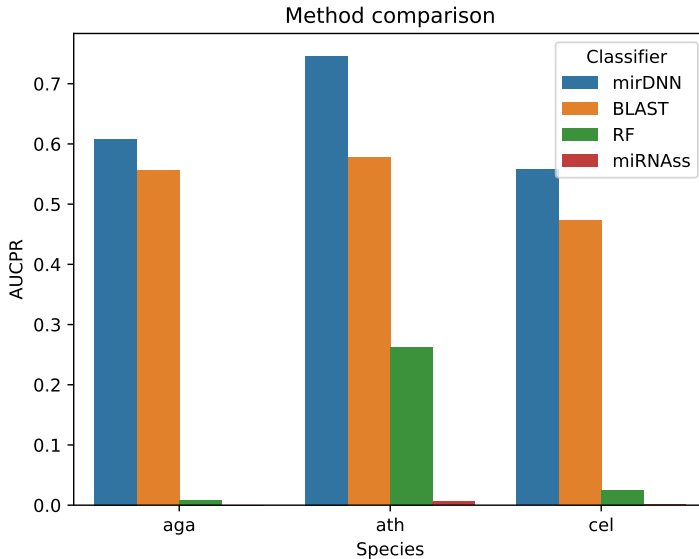
Precision-recall curves



Precision-recall curves



Area under the curves



- 1 Introduction
 - Background
 - Motivation
- 2 Proposed method
 - Convolutional Neural Networks
 - Proposed architecture
- 3 Results
 - Experimental setup
 - Precision-recall curves
- 4 Conclusions

Conclusions

- State-of-the-art machine learning methods for miRNA prediction do not perform well on new species.
- Sequence alignment methods have a better performance, but only for low recall rates.

Conclusions

- State-of-the-art machine learning methods for miRNA prediction do not perform well on new species.
- Sequence alignment methods have a better performance, but only for low recall rates.
- Using convolutional neural networks, we can *learn* which are the important characteristics that define a pre-miRNA.

Conclusions

- State-of-the-art machine learning methods for miRNA prediction do not perform well on new species.
- Sequence alignment methods have a better performance, but only for low recall rates.
- Using convolutional neural networks, we can *learn* which are the important characteristics that define a pre-miRNA.
- The proposed method achieves a precision many times higher at high recall rates in three tested species.

Conclusions

- State-of-the-art machine learning methods for miRNA prediction do not perform well on new species.
- Sequence alignment methods have a better performance, but only for low recall rates.
- Using convolutional neural networks, we can *learn* which are the important characteristics that define a pre-miRNA.
- The proposed method achieves a precision many times higher at high recall rates in three tested species.

MicroRNA prediction from genome-wide data with deep learning: a novel approach based on convolutional residual networks

C. Yones, L.A. Bugnon, J. Raad, D.H. Milone and G. Stegmayer

Research Institute for Signals, Systems and Computational Intelligence



`cyones@sinc.unl.edu.ar`
`http://www.sinc.unl.edu.ar/`